

COUNTERFACTUAL AFFORDANCE MAPS: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study robot affordance learning. The motivating bet is: Map affordances that would exist under alternate grasps, poses, and supports.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For robot affordance learning, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2025a;b; 2024; 2021; 2023; 2025c). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.512	0.488
observed_only_wam	hidden_mode_shift	0.193	0.807
augmentation_as_data_wam	nominal	0.573	0.427
augmentation_as_data_wam	hidden_mode_shift	0.282	0.718
scalar_uncertainty_wam	nominal	0.685	0.315
scalar_uncertainty_wam	hidden_mode_shift	0.370	0.630
proposed_branch_mechanism	nominal	0.802	0.198
proposed_branch_mechanism	hidden_mode_shift	0.535	0.465

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Counterfactual Affordance Maps argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Affordance-based grasping point detection using graph convolutional networks for industrial bin-picking applications, 2021. <https://doi.org/10.3390/s21030816>.
- Non-standard map robot path planning approach based on ant colony algorithms, 2023. <https://doi.org/10.3390/s23177502>.
- Gap-rl: Grasps as points for rl towards dynamic object grasping, 2024. <http://arxiv.org/abs/2410.03509v1>.
- Gat-grasp: Gesture-driven affordance transfer for task-aware robotic grasping, 2025a. <http://arxiv.org/abs/2503.06227v1>.
- Sent map – semantically enhanced topological maps with foundation models, 2025b. <http://arxiv.org/abs/2511.03165v1>.
- Map-vla: Memory-augmented prompting for vision-language-action model in robotic manipulation, 2025c. <http://arxiv.org/abs/2511.09516v1>.